

Practicals for Data Analysis & Programming for Operations Management

N.D. Van Foreest, W. van Wezel

1. INTRODUCTION

This document contains the material for the practicals related to optimization of Operations Management problems with Python, Gurobi, and Elasticsearch. We rely on you to search on web for examples, documentation, and so on. For instance, a search on Python lp examples will provide you with lots of useful websites that provide code examples on how to use python to solve LP problems. In fact, by reading (thoroughly!) code examples of others, we learned to program in Python ourselves. First we looked for simple examples, then we read harder examples, and so on.

At the end of each practical week we will make our code available, so that you can compare your solution with ours. Given that students will share their code anyway via, e.g., Dropbox anyway, there is no point in not giving you the code. In fact, by giving you our code, we can at least ensure you get good code. However, please try real hard to complete the practicals yourself.

2. PRACTICAL 1B: PRODUCT-MIX OPTIMIZATION

2.1 *Get a working python environment with Gurobi*

For the practicals, you will have to install a working Python environment and the mathematical optimization solver from Gurobi, which is perhaps the best LP optimization library available, used by nearly every multi-national with optimization problems such as KLM, Schiphol, IBM, and so on). See <https://www.gurobi.com/>. How to get the Python environment is explained in the manual for the a-practicals (generic programming skills). So here, we assume you have a Python and Atom installed on your computer. To install Gurobi and link it to Python, follow these steps:

1. Download Gurobi on <https://www.gurobi.com/>. On the top of the screen, click 'downloads & licenses', and then click on 'download center'.
2. In the download center, click on 'Gurobi Optimizer'
3. Accept the license
4. Then download the appropriate version of the latest release of the Gurobi Optimizer (this depends whether you are on Windows, MacOS, or Linux).
5. You can then install Gurobi.
6. Now go back to the download center, and, in the 'request a license' section, click on 'Academic License'.
7. You need to register an account. Make sure to use your university email address; they might require this to verify you are indeed a student from academia.
8. When you have registered, you can generate a license code.
9. When generated, you can activate this license code on your computer:
 - a. Open a command prompt (Windows: click the windows button, type cmd, and press enter; on MacOS: open a terminal).
 - b. Copy the command provided by the Gurobi website and press enter. The command should look like this (the numbers will be different for you):
`grbgetkey 123456-abcde-321232-6456-2342432432`

If you don't get error messages, Gurobi is now installed. We now must arrange for Python to be able to use Gurobi. This requires the installation of the Gurobi Python libraries. The easiest way to do this is start a command prompt or terminal again, and go to the directory where Gurobi is installed. For me this was C:\gurobi903\win64. So after opening the command prompt, I type `cd C:\gurobi903\win64`, and then press enter. This folder should contain a file called `setup.py`. On your command prompt, now type `python setup.py install`. This will install the Gurobi libraries in Python.

2.2 *Running the product mix code*

1. Download `product_mix.py` from nestor; the code is also available in Table 1. This code implements a linear program (LP), see Factory Physics Ch 16 for further explanation.
2. Run this code in Atom.
3. Tip: if you run the code and you get the output in a small popup window, then right-click on the code window, click 'hydrogen' in the menu, and then click 'toggle output area'. Now run the code again. The output will now be in the right window.
4. Read and think about the output.

2.3 Breaking the code

It is extremely useful to become familiar with python's error message. When you are coding, you'll often make mistakes (this is entirely normal; it happens to us all the time). Python's error messages can be enormously helpful in discovering what you did wrong. For this reason, we are going to break our working code, and see what type of errors we get. Read the error trace carefully; for instance, check the line number where the error occurs, and check the error type.

Here are some suggestions to break the code. Do them all, and then try your own mistakes. Read the error message very carefully. Understanding the error messages can save you an enormous amount of time (not just minutes or hours, but days!) After introducing the error, don't forget to repair, i.e., undo, it.

1. Comment Line 10, see Table 1, in your code; that is, put a `#` in front of the line. (Mind that python is sensitive to indentation, hence, the `#` should be the very first character on the line, there should be NOTHING in front of it). What is the error message? Remove the comment sign again.
2. Now comment out Line 12 (and uncomment it after having read the error message).
3. Line 12, put a space in front of the first word, so that everything moves one step to the right. Check the error message, and explain it.
4. Line 12, what would happen if you would put two spaces at the front of the line?
5. Line 12, what would happen if you would put a tab character at the front of the line?
6. Comment out Line 14.
7. In Line 14, change `addVar` to `addvar`. (Just change the capital 'V' to 'v').
8. In Line 14, change `x1` to `x`.
9. Comment out Line 17.
10. Comment out Line 19.
11. Comment out Line 21.
12. Line 21, remove the number 2400 so that the line reads like: `m.addConstr(15 * x1 + 10 * x2 <=)`

13. Comment out Line 26.
14. Comment out Line 29. What happens? Explain it.
15. Line 26: replace the line by `for v in:`
16. Introduce your own errors, one by one, and repair again. The more errors you can invent, the better. Once again, reading and understanding python's error messages will save you lots and lots of time, and frustration.

Table 1: product_mix.py

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

```

1  #!/usr/bin/python
2
3  # This example implements the product mix example of FP,
4  # Appendix 16.A. See eqs 16.107--16.113.
5
6  # Nicky van Foreest, 2019
7
8  %reset -f
9
10 from gurobipy import Model, GRB
11
12 m = Model("product mix")
13
14 x1 = m.addVar(ub=100, name="x1")
15 x2 = m.addVar(ub=50, name="x2")
16
17 m.setObjective(45 * x1 + 60 * x2, GRB.MAXIMIZE)
18
19 m.addConstr(15 * x1 + 10 * x2 <= 2400)
20 m.addConstr(15 * x1 + 35 * x2 <= 2400)
21 m.addConstr(15 * x1 + 5 * x2 <= 2400)
22 m.addConstr(25 * x1 + 14 * x2 <= 2400)
23
24 m.optimize()
25
26 for v in m.getVars():
27     print("%s %g" % (v.varName, v.x))
28
29 print("Obj: %g" % m.objVal)
30

```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

2.4 Integer variables, get familiar with the gurobi documentation.

In the solution you must have seen that the optimal number of products to be produces is a fraction. Check the Gurobi documentation how to enforce that x_1 and x_2 are integer. (Just search the web. . .) Hint, google for `gurobi addvar integer`.

You should know that optimization problems whose variables are constrained to integers are typically much, much more complicated than problems all of whose variables are

continuous. If discrete variables are not necessary, do not include the GRB.INTEGER property! Including this property in the problem specification for Gurobi, i.e., including the integer property in the code, can easily change the computation time from seconds to years for industry size problem. To understand the problem, suppose that the optimal value of x_1 is either 75 or 76, and for x_2 it is either 35 or 36. To find the optimal integer solution we need to test $2 \times 2 = 2^2 = 4$ alternatives. If we have 100 decisions variables, we need to test about 2^{100} possibilities; this is a huge number, impossible to test in any reasonable amount of time. So be aware!

2.5 Finding the constraints

Ensure that your code is equal to the original code again, so that you have a correctly working example. Let Gurobi solve the problem. Then you should notice that not all demand is covered. By changing the production system it must be possible to increase sales. But. . . , which machine to change, and what to do? In the next couple of exercises we will tackle this problem.

You should observe, and remember, from these problems how easy the numerical analysis becomes with our Python/Gurobi environment. We generate some ideas to extend production capacity. To see the effect of these suggestions, we have to change some numbers in our code, but that is very easy. Then we let Gurobi solve the problem, and we read out the results, and compare the effect on the total revenue. In other words, we (as humans) suggest plans on how to improve things, we let the computer carry out all the boring computations, and then we (as humans) interpret the results.

Now, our plan is to add capacity. But which machine(s) should we update? Which is the best to invest in?

A simple (but rather dumb) way to find out the machine that is the most constraining is as follows. We remove a machine from the set of constraints in lines 16–19, and just check whether the output (i.e., the value of the objective) increases.

1. First comment Line 19, and let gurobi solve the problem. Check what happens to the solution.
2. Now include Line 19 again, but comment Line 20. Let gurobi solve the problem and check the solution.
3. And so on.

You should see that machines B and D have an impact on the revenue.

You should know that the above procedure is by far not the smartest. You can use Gurobi to identify the, so-called, active constraints. However, this requires more knowledge of Gurobi, which you can always acquire later.

2.6 *Extra minutes for the constraining machines*

The learning goal for this and the next session is that you understand how easy it is with computers to quantitatively analyze many different scenarios. As an example, suppose we can buy some overtime, let's say 100 minutes per week on any machine. We would like to know which machine is the best. The next couple of questions are meant to help you get started with asking and answering such questions.

1. Change the constraint on machine B to 2500. Then solve the problem again, with Gurobi. And check the result.
2. What happens if you just add 1000 minutes to machine B? Realize that some of this extra time might be unused. Then this is not-so-smart decision: we pay for personnel, but we are unable to exploit this extra capacity due to the other machines.
3. Undo the changes on machine B, but change machine D's time to 2500. Then solve the problem again, with Gurobi. And check the result.
4. Perhaps it's better to give 50 minutes to machine B and D. What is the result?
5. How many extra minutes are necessary to satisfy all demand? (You don't have to find a general procedure for this problem, just play a bit with the numbers to see the effect.)
6. Another idea is to move some of the total time of machine A to machine B. In other words, we train the personnel that operates machine A to help at machine B, thereby decreasing the total available time of machine A (but reducing the time of machine A). Let's assume that the operators of machine A are less fast when they help at machine B. For instance, if we remove 200 minutes from machine A then this can be converted to just 100 minutes at machine B. Should we consider training machine A's personnel to help at machine B? Finding out is easy: change the 2400 of machine A into 2300, and increase the 2400 to 2450 at machine B. Have Gurobi solve this and interpret the result.
7. Make a graph of the output as a function of the extra capacity added to machine B. For the present, it is ok to make this in Excel. Later you'll learn how to make this with Python tools.

We assume that by now you understand how easy scenario analysis becomes. You should realize that scenario analysis is one of the key skills of managers. Generate good ideas on what to change, analyze the effects of the changes, interpret the results, and implement the best change.

2.7 *Moving work from a bottleneck to a non-bottleneck*

Another way to shift work from machine B to machine A (or C) is to reduce the production time. The idea is to assume that part of the work of an item can be done at either machine.

1. (Don't forget to bring your code to a pristine state again.)
2. What is the impact of increasing the production time of product 1 at machine A from 15 to 20 minutes, but thereby reducing the production time at station B from 15 to 10 minutes?
3. Should we move time for product 1 to station A or station C?
4. or should we move production time of product 2?
5. Should we move time from station D to station C? Which product?
6. How much time should we try to move? This question is particularly interesting. Suppose we would be able to redesign product 1 so that we can move all its production time at machine B to machine A, so that machine B can be bypassed altogether. Would it be useful to investigate such a design? (If such a redesign does not result in much extra output, it's useless to try. So, with our computations we can first analyze the effect on logistics and revenue before we start thinking about how to do the redesign.)

Assignment Project Exam Help

2.8 Testing

Finally, we need to check whether our model is correct. Note that for a real case it is essential to start with testing your model and implementation. In this practical we did the scenario first, so we sinned a bit. The reason is to show you first how to use computational tools such as Python and Gurobi, and not bore you with testing. But remember, when you use tools to make decisions for real, then always test first. It's easy to make a typo (e.g., type in the wrong number), miss a constraint, and so on, and it is a bit painful when your multi-million decision turns out not to work because you typed in the wrong numbers. . .

1. Remove all machine constraints, by commenting them out, and check that only the demands form the constraints.
2. Include the machine constraints again, but now remove the machine constraints on the demands, by removing the upper bounds indicated by the key word 'ub' in the code. (Check Gurobi's documentation on how to remove the upper bound.) Use Figure FP 16.15 to see what would be the optimal solution for this case, and check that you get the same solution with gurobi.

2.9 Adding Machines

Yet another idea is to buy an extra machine that can take off some of the load of machine D, say. So, we are going to add a machine E to the end of the chain of machines.

1. Copy `product_mix.py` to `product_mix_fifth_machine.py`, and use the latter file to implement the changes.
2. Add to the LP a machine E with production times 13 and 7 for products 1 and 2, respectively. Machine E is available for 1500 minutes per week. The production times at station D change now to 12 and 10 respectively.
3. Analyze the quantitative effect of moving work from machine D to machine E. (I see an increase in revenue from 5575 to 6042 by adding a machine E.)

So, based on estimates of the cost of buying and operating such a machine E we can use our computer programs to make good decisions, i.e., whether it's worth the money or not.

2.10 Adding a third product

A marketer suggests to add a third product to our product portfolio. The estimated production times are 10, 5, 18, 70 on the respective machines; the selling price is \$120 and the cost of raw material including labor is \$30. The demand is estimated to be 20 products per week. Should we include this product in our portfolio?

1. Copy `product_mix.py` to `product_mix_third_product.py`.
2. Add product 3 to the LP and analyze the effect, in particular on the total profit. Does the profit increase when you add product 3?
3. Due to contractual obligations we have to minimally produce 80 items of product 1 and 20 of product 2. Include this in the LP. (tip: in `addVar`, `ub` stands for upper bound. Check the impact.

2.11 Minimizing the capacity required

Here is an interesting challenge: What would be the minimal capacity per machine required to serve all demand? You can build this also as an LP. Here I leave the details to you. The solution becomes available at the end of the week. Tip: what should be your decision variables now?

Upload your code of 2.11 to Nestor as end product of Week 1.

3 PRACTICAL 2B: CODE ORGANIZATION AND INVENTORY CONTROL

The code in Table 1 has some problems. It does not scale to large numbers of machines or products because each product and machine has to be included by 'hand'. Moreover, the data is hard-coded in the algorithm itself. It is a much better design pattern to separate data from the algorithm, and make the algorithm generic in the sense that it can scale up to many products or machines. That is what we are going to learn in Section 3.1.

In Section 3.2 we'll organize our code a bit more with functions. With this step we complete the product mix example.

Then, in Section 3.3, we will analyze an inventory control problem with LP. This is a really useful example; over the years students have been applying it at companies to control the production and inventories.

In the next practicals we will no longer take you by the hand, but just give you hints on how to work towards an interesting computer program.

Assignment Project Exam Help

3.1 *Separating data from the formal problem specification*

We continue extending the case of the previous practical.

1. Create a new file, `product_mix_2.py`, so that you don't mess up the program in `product_mix.py` (This one worked, so while developing something new, it is best not to break the stuff that works.)
2. Put demand data in a list, called array `D`.
3. Put profits in a list `P`.
4. Put the capacity constraints of the machines in a list `C`.
5. Put the production times in a list of lists, `PT`, such that, for instance, `PT[0][1]` corresponds to the production time of product 2 at machine A (recall that python starts counting at 0, hence we subtract 1 from the machine id and product id).
6. Next, we need to make decision variables. Search, for instance, on `gurobi multiple variables example`. Figure out how to make a list (also called array) of decision variables.
7. We now need to add the constraints. A single constraint often refers to multiple decision variables and data elements. In our case, we need to multiply the decision variable with the processing time, and compare the result to the maximum machine capacity. In the previous version of the code, we manually added a constraint for each machine. Now that the machines and products are configurable we need to add the constraints in a loop. Conceptually, we would need something like:

```
for i in range(len(C)):
    m.addConstr(time_used_on_this_machine<=C[i])
```

where the time used on the machine would be the multiplication of all products produced on that machine multiplied by the amount produced (which is our decision variable x). So, within a single constraint, we need to iterate over all products. This is where generators and the Python `sum` function can help us. For example, try the following code:

```
x = sum(2 * n for n in range(5))
print(x)
```

The code within the `sum` function is called a generator. Track that this code is a shorter version of:

```
x = 0
for n in range(5):
    x = x + 2 * n
print(x)
```

(ps, if you expected the outcome of this code to be 30, please read https://www.w3schools.com/python/ref_func_range.asp)

This can also be used when adding constraints. It could look something like:

```
num_machines = len(C)
for i in range(num_machines):
    m.addConstr(sum(PT[i][j] * x[j] for j in range(num_products)) <= C[i])
```

But this formulation contains an error. Repair the error, and modify it so that you can use it in your code.

8. You also need the `sum` function in your `setObjective` call. Complete your program until it runs. Then compare the outcome with your original `product_mix.py`. The results should match, otherwise there is a mistake somewhere.

3.2 Organize code in functions

The challenge here is to organize the code in a better way, that is, by means of functions. Functions are very useful for a number of reasons. The first is that function names (if properly chosen) act as documentation; a good function name describes what the code in the

body does. The function also helps to hide complexity. A programmer just has to read the function name to guess what the function does; the programmer does not necessarily have to read all the code of the function. Finally, a function enables reuse; it is not necessary to copy the code time and again for nearly the same task but with different numbers; just call the function again with new arguments/parameters.

1. Make a new file `product_mix_3.py`.
2. Give this script the same functionality as the previous one, but now make a function `optimize` to which you pass the data arrays as argument. The output of the function should be the optimal profit.
3. Run it to see that it works.

3.3 Inventory control

We will now implement the model of Section 16.2.1 of FP. As said, it is an interesting and practically useful case; students have been using it as part of their master thesis project to control the production and inventory levels at companies. You should realize that in industry problems, there can be like 1000 different product types, the holding costs and profits may depend on product type and time (e.g., discounts in summer time), the inventory positions run in the millions of items. Such cases are impossible to analyze with Excel; programming skills are absolutely necessary here¹.

1. Read FP Section 16.2.1 and the one paragraph intro of Section 16. We use the data of Section 16.2 to check our implementation.
2. Make a list with demand and capacities like so


```
D = [0, 80, 100, 120, 140, 90, 140]
C = [0, 100, 100, 100, 120, 120, 120]
```

 There is a reason for this that is slightly subtle. If you check FP, Eq 16.4, you'll see that to update the inventory level I_1 at time $t = 1$, you'll need the inventory level I_0 at time $t = 0$. Thus, in the implementation, the easiest is let all lists start at time 0 and put zeros where appropriate.
3. Add decision variables for X , I , and S . Note that they need to be 7 long, 6 for the number of periods $t = 1, \dots, 6$, and 1 for period $t = 0$. Include in the definition of X , S and I that they should be non-negative, i.e., FP Eq. 16.5. Then include also that $X \leq C$ and $S \leq D$, i.e., Eqs 16.2 and 16.3.
4. Add constraints to ensure that $X[0] = 0$, $S[0] = 0$, and $I[0] = I_0$. In the specification of the problem in Section 16.2 $I_0 = 0$.

¹ In fact, as a general rule, the analysis of any realistic business problem cannot be carried out with Excel, or it is exceptionally hard (much, much harder than to write a python program and use libraries). So, once again, learn to program.

5. Finally add constraint 16.4 for every t with a for loop. Observe that we need 6 constraints in total. Observe also that if you do have to deal with weekly or daily demand data, for one year, than adding all constraints by typing is not particularly nice. You really need for loops here.
6. Finally, add objective Fp. Eq. 16.1; use `sum` or `quicksum` to make your life easy. Quicksum is a Gurobi specific version of `sum`, which is faster for optimization models. If you use it, make sure to also import the `quicksum` library.
7. Solve the problem, and compare your result with the output as shown in Figure 16.5.
8. Think a bit about how your excel sheet would look like if you would had to deal with yearly data. You should see that it would not fit on screen; in fact, it would look extremely messy.²
9. Write a function to read demand and capacity data from the file `inv_control_data_1.txt`. (First open this file with, e.g., Atom or notepad. It is best to check the format of an input file before you write code to read it.) Each line contains the demand and the capacity for a specific period. The data is the same as earlier, so testing must be easy. Note that if you read the file as csv and convert it to a list, the individual items in the list will be a string. You need to tell Python explicitly that the values are numbers, for example, `int(record[0])` if you want the values to be rounded to a whole number or `float(record[0])` if the numbers could contain fractions (e.g., 3.8).
10. Then read the info from `inv_control_data_2.txt`; this represents one year of weekly data. Solve the optimization problem. Interpret the results. When, for instance, does inventory start to build up? Are there weeks that we have lost sales? If so, is r sufficiently large? Can you modify the capacity planning to maximize the profit? (This is interesting challenge, and in fact, the challenge in many master thesis projects.)
11. An interesting optional extension is to try to make a graph of the inventory as a function of time. Graphs can be made with the `matplotlib` library.

Upload your inventory control model code to Nestor as end product of Week 2.

² As a general lesson again, Excel files are very hard to maintain and understand in comparison to (Python) programs. My code for this problem consists of some 15 lines, each of which is completely clear by itself.

4 PRACTICAL 3B: MAKE YOUR OWN STRAVA

In this practical you'll build your own sort of Strava tool (if you don't know what it is, see <https://www.strava.com/>). Of course, it will not be as slick as the real Strava. However, you don't get unwanted advertisements; you can process your own data, and, on top of this, you'll come to understand how data assembly of routing works.

Although Strava is not really useful in a business context, geo-location tools are very important in transportation companies. Not only for traditional transportation planning and routing, but even more so for new logistic business concepts such as Uber, thuisbezorgd.nl, and the Physical Internet. Now that gps data of all trucks and couriers can be shared real time with planners, such companies get their main competitive advantage from dynamic planning using real time locations of couriers and real time data of traffic congestion. Therefore, understanding and being able to work with geo-location data and tools is very important for a data scientist.

4.1 *Computing track length and duration*

In our first version of our tracking tool we'll use a library to read the data from a gpx file and we use a library to compute the distance between two gps coordinates. The goal is to build a computer program that computes the total length and duration of a track recorded by an app that runs on your mobile phone.

1. Install an app that can track your routes and can save the data in the standard gpx format. For my iphone I use the app 'Open GPX Tracker'; for Android there must be something similar.
2. Record a route and send it to yourself by mail. If you don't want to track your own route, use the gpx file available on Nestor. Open the file in Atom and look at the structure. It is formatted as Xml, which is a hierarchical structured file format, primarily used to exchange information between applications. You see it has `lat`, `lon`, and `time` tags. Apparently, the GPX tracker has recorded the gps location every few seconds (see https://en.wikipedia.org/wiki/Geographic_coordinate_system).
3. Consider that, using the string parsing functions in Python, you could extract this data yourself. However, others have already created this functionality, so if something exists that is already used by many others and extensively tested, it is better to use that. Therefore, install the `gpxpy` package to read the info from the gpx file. Remember, you can install packages using `pip` at the command prompt

- or terminal: `python -m pip install gpncpy`. (On the Mac, you need to use the python version that is also used by Atom, e.g., `python3.8` or `python3.7`).
4. The code examples at <https://github.com/tkrajina/gpncpy> show you right away how to use the gpncpy package. Copy the code and print for instance the latitude of all points from the gpx file.
 5. Once you can print the latitudes of all points, you know you are on the right 'track'. Now remove the print, but try store the points as a list of points (call this list `points`), such that each point in the list is a list itself with the three items `[latitude, longitude, time]`.
 6. Once you have the list of points, you should realize that `points[-1]` is the last point, and `points[0]` the first. Now explain that `duration = points[-1][2] - points[0][2]` is the total duration of the track.
 7. It is important to know that duration is a `timedelta` object. With `duration.seconds` you'll get the duration in seconds. With these ideas, write code to print the duration of the track in seconds. Open the gpx file in Atom (or some other editor) to check whether your result makes sense. Just check the time stamp of the first point in the gpx file and the time stamp of the last point.
 8. The next step is to compute the total distance traveled. For this, use the `geopy` package, or some similar package to compute the distance between points with gpx coordinates. (Once again, I did not want to find out how to compute such distances, I just know somebody has solved this, so my problem reduces to finding the right package that does the job for me.)
 9. Compute the distance between the first two points as follows. Put `from geopy import distance` at the top of your program, so that you can use the distance object from the geopy module. Then use the following code to compute the distance:

```
p1, p2 = points[:2]
print(
    "distance between first two points in meters: ",
    distance.distance((p1[0], p1[1]), (p2[0], p2[1])).meters,
)
```
 10. Once you know how to compute the distance between the first two gps points , you should find a way to add all distances together to compute the total distance.
 11. Finally, now that you have the duration and the total distance, it's easy to compute the average speed. Complete the code and print all results.

So, now we have the raw data; what can we do with it? We could for instance compute the speeds at each time point, and then find the maximal speed. This turns out to be a bit tricky, so have we to modify the analysis.

The lesson of the section is that, in data science, the transformation of raw data to something that is useful and realistic (i.e., fits in the context), nearly always requires quite a bit of additional effort. In other words, the analysis starts, not ends, with the raw data.

1. Iterate through all segments, and calculate the speed per segment.
2. Compute the minimal and maximal speed; check the internet to how to efficiently find the maximum element in a list. The highest speed in the gpx file from Nestor should be around 10 km/h.

If you would have used the browser to plot our track on a map, you would have seen that it is a walk around the Duisenberg building. In fact, we recorded it for this course when we walked around the building. Certainly we did not run, so the max speed of about 10 km/h is impossible.

1. One guess is that the time measurements are somewhat irregular. With the above tools you can make an algorithm to compute the time differences between the successive points and then compute the largest and smallest difference. Is there a large spread in these time differences or not?
2. Can you find a reason on the web about why this problem (of too high speeds) occurs?
3. As a challenge, can you find a method to repair for this problem? (Hint, the technique is called ‘smoothing’; it quickly becomes quite mathematical, and also very interesting.)
4. Optionally, you can watch the following talk about a more advanced way to correct for errors in gps measurement:
https://www.youtube.com/watch?v=9Q8nEA_0ccg

In this case we have to repair outliers in the speed measurements. Errors in raw data occur always. Some examples are measuring the temperature of a machine to predict a breakdown, website usage statistics (which could be polluted by visits of search engines), and typo’s when machine operators must enter data into a system. It is often difficult or impossible to determine whether an outlier is a measurement error or a real outlier. In general, you should make a number of different measurements (such as in our case the max, mean, and min of the speed) to obtain an impression of the quality of your data. Then you should do some (or a lot of) massaging of the raw data to get insights you need.

4.3 *Plotting the route*

The gps coordinates themselves are of little values for us. Luckily, we can easily plot the route on a map.

1. For this, you can use the `smopy` package, which can be installed with `pip`.
2. Study the example code on <https://github.com/rossant/smopy>.
3. When initializing the map, you need to specify the approximate part of the map you wish to show. For this, lookup the minimum and maximum latitude and longitude using your program, and use these as parameters when you initialize the map.
4. Another parameter in the map initialization is the zoom level. The website mentions `z=4`. You can increase this depending on the route you walked. With the gpx file on Nestor, a zoom level of 14 is sufficient.
5. If you wish to show the map a bit bigger, increase the size in the `show_mpl` call.
6. By plotting the individual gps locations as points on the map, you can see the route. To test this functionality, first try to show the first gps location in the gpx file.
7. If this works correctly, iterate and plot all points on the map. If you think the points are too thick, use `ms=3` and `mew=1`.

Upload the map plotting code to Nestor as end product of Week 3.

4.4 Optional extensions

1. A plot of the speeds in a graph (as determined in 4.2) is quite interesting. For instance, Uber eats can figure out how long payments typically take. When the courier is at the door of a client, the speed must be zero.
2. You can determine the color of a plotted point through code, for example:

```
ax.plot(x, y, 'or', color="green", ms=3, mew=1);
```

This can be used to indicate the speed per segment on the map. Try to give different colors for segments where the speed is below 4, between 4 and 7, and above 7 km/h.
3. Other applications are very easy now too. If you want to know your total amount of kilometers walked in a month, just track your walks. Put the tracks that are obtained with walking into a specific directory (to prevent messing up your walking data with your cycling data, for instance). Then, for the total distance, just add the total distances per track (after you corrected the data for suspect measurements of course).